

Report - Quora Question Papers

Problem Statement

The problem statement is to predict which of the provided pairs of questions contain two questions with the same meaning.

Data

Our task is to build classification models which will classify the target variable(is_duplicate)

id - the id of a training set question pair

qid1, qid2 - unique ids of each question (only available in train.csv)

question1, question2 - the full text of each question

is_duplicate - the target variable, set to 1 if question1 and question2 have essentially the same meaning, and 0 otherwise.

Structure of Data

```
> str(data)
'data.frame': 404290 obs. of 6 variables:
 $ id      : int  0 1 2 3 4 5 6 7 8 9 ...
 $ qid1    : int  1 3 5 7 9 11 13 15 17 19 ...
 $ qid2    : int  2 4 6 8 10 12 14 16 18 20 ...
 $ question1 : chr  "What is the step by step guide to invest in share market in india?" "What is the story of Kohinoor (Koh-i-
the speed of my internet connection while using a VPN?" "Why am I mentally very lonely? How can I solve it?" ...
 $ question2  : chr  "What is the step by step guide to invest in share market?" "What would happen if the Indian government sto
back?" "How can Internet speed be increased by hacking through DNS?" "Find the remainder when  $23^{24}$  is divided by
 $ is_duplicate: int  0 0 0 0 0 1 0 1 0 0 ...
```

Summary of Data

```
> summary(data)
      id      qid1      qid2      question1
Min.   :    0   Min.   :    1   Min.   :    2   Length:404290
1st Qu.:101072 1st Qu.: 74438 1st Qu.: 74727   Class :character
Median :202145 Median :192182 Median :197052   Mode  :character
Mean   :202145 Mean   :217244 Mean   :220956
3rd Qu.:303217 3rd Qu.:346574 3rd Qu.:354693
Max.   :404289 Max.   :537932 Max.   :537933
      question2      is_duplicate
Length:404290   Min.   :0.0000
Class :character 1st Qu.:0.0000
Mode  :character Median :0.0000
                  Mean   :0.3692
                  3rd Qu.:1.0000
                  Max.   :1.0000
.
```

Methodology

As per the problem statement we need to analyze text which requires different techniques as compared to numerical or categorical data.

We can approach this problem using various text processing methods which can give the meaning of the text in numerical terms which makes it easy for us to analyze this data.

Stringdist package is extensively used in this project. In which different methods of string comparisons are used.

Let's start

```
> length(unique(data$question1))
[1] 290457
> length(unique(data$question2))
[1] 299175
```

The data contains many repeated questions

Count of Duplicate levels

```
> table(data$is_duplicate)

      0      1
255025 149263

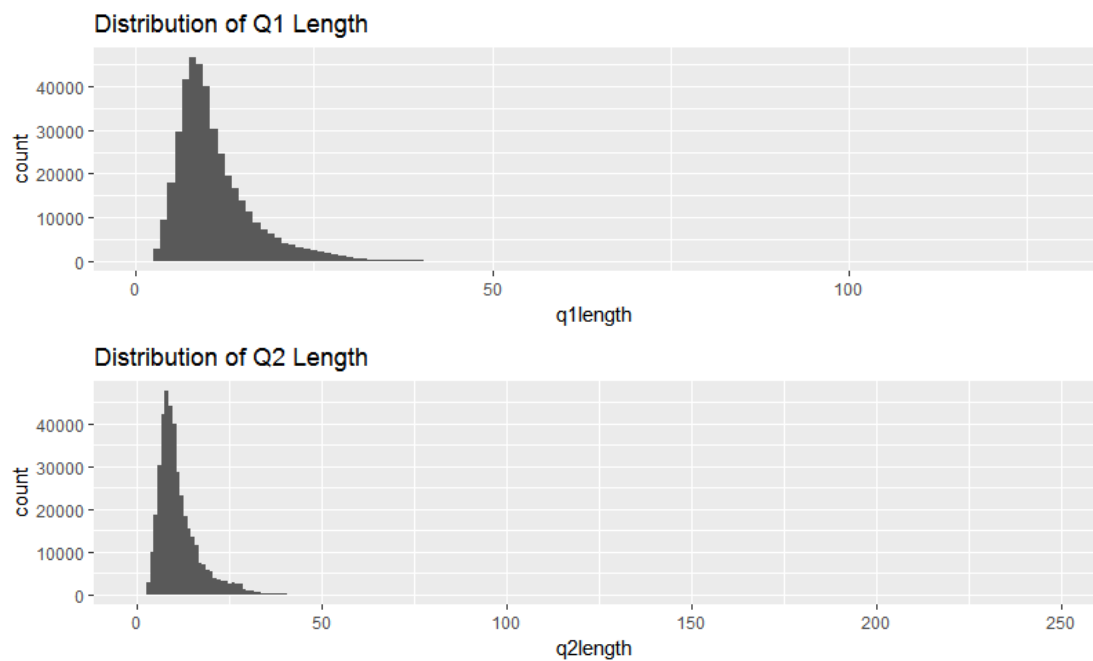
> sum(data$is_duplicate==0)/nrow(data)*100
[1] 63.08021
> sum(data$is_duplicate==1)/nrow(data)*100
[1] 36.91979
```

We have 63% of the observation of class 0 in our data.

Distribution of length in question1 & question 2

```
> summary(text_process[,4:5])
      q1length      q2length
Min.   : 1.00   Min.   : 1.00
1st Qu.: 7.00   1st Qu.: 7.00
Median :10.00   Median :10.00
Mean   :11.13   Mean   :11.38
3rd Qu.:13.00   3rd Qu.:13.00
Max.   :128.00  Max.   :247.00
```


From the summary we see that the average length is around 7-13 but some question show out of the ordinary lengths which can be seen in below box plot



```
> sum(text_process$q1length>=7 & text_process$q1length<=13)/nrow(text_process)*100
[1] 61.27626
> sum(text_process$q1length<7)/nrow(text_process)*100
[1] 14.83474
> sum(text_process$q1length>13)/nrow(text_process)*100
[1] 23.889
>
> sum(text_process$q2length>=7 & text_process$q2length<=13)/nrow(text_process)*100
[1] 60.26259
> sum(text_process$q2length<7)/nrow(text_process)*100
[1] 15.23249
> sum(text_process$q2length>13)/nrow(text_process)*100
[1] 24.50491
```


Around 60% question length falls between 7 & 13.

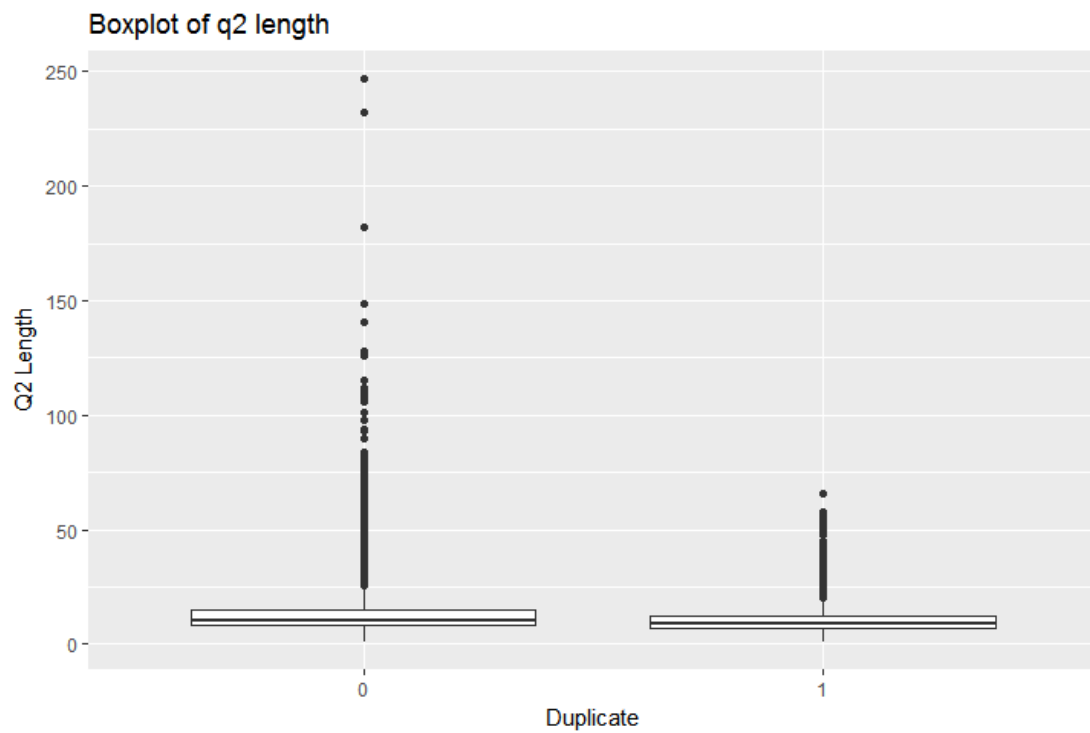
Outlier Analysis on text length for two levels of Duplicate




Q1 & Q2 length boxplot

Q1 & Q2 Box Plot for each level of Duplicate









Density Plot for Q1 & Q2 Length



Let's see it with bar plot



This shows as the length increases difference in duplicity becomes apparent.

Pre-processing

Data pre-processing is done through **regular expressions** all punctuation marks are removed, text is converted to their base forms, text is converted to untext code . Numbers are not removed from the data as they may hold significance.

Let's do our work on text analysis with stringdist package which provides us many methods to work on text.

1. **Q-gram**

It is a subsequence of q consecutive characters of a string. If x (y) is the vector of counts of q -gram occurrences in a (b), the q -gram distance is given by the sum over the absolute differences $|x_i - y_i|$.

2. **Jaro Distance**

The Jaro distance (method='jw', p=0), is a number between 0 (exact match) and 1 (completely dissimilar) measuring dissimilarity between strings. It is defined to be 0 when both strings have length 0, and 1 when there are no character matches between a and b .

3. **Cosine Distance**

Cosine similarity is a measure of similarity between two non-zero vectors of an inner product space that measures the cosine of the angle between them.

4. **Stringsim**

It computes pairwise string similarities between elements of character vectors a and b , where the vector with less elements is recycled. The similarity is calculated by first calculating the distance using stringdist, dividing the distance by the maximum possible distance, and subtracting the result from 1. This results in a score between 0 and 1, with 1 corresponding to complete similarity and 0 to complete dissimilarity

5. **Longest Common Substring (LCS)**

The longest string that can be obtained by pairing characters from a and b while keeping the order of characters intact. The lcs-distance is defined as the number of unpaired characters. The distance is equivalent to the edit distance allowing only deletions and insertions, each with weight one.

6. **Full Damerau-Levenshtein distance**

It is like the optimal string alignment distance except that it allows for multiple edits on substrings.

7. **Jaccard Distance Method**

It is a statistic used for comparing the similarity and diversity of sample sets.

8. **Optimal String Alignment distance (OSA)**

It counts the number of deletions, insertions and substitutions necessary to turn b into a . allows transposition of adjacent characters. Here, each substring may be edited only once.

After applying all the methods we have our data which is quantitative and this allows us to perform numerous algorithms.
Let's look at our data now

```
> str(text_process)
'data.frame': 404274 obs. of 14 variables:
 $ question1 : chr "what is the step by step guide to invest in share
market in india" "what is the story of kohinoor koh i noor diamond" "how
can i increase the speed of my internet connection while using a vpn" "wh
y am i mentally very lonely how can i solve it" ...
 $ question2 : chr "what is the step by step guide to invest in share
market" "what would happen if the indian government stole the kohinoor ko
h i noor diamond back" "how can internet speed be increased by hacking th
rough dns" "find the remainder when math 23 24 math is divided by 24 23"
...
 $ is_duplicate: Factor w/ 2 levels "0","1": 1 1 1 1 1 2 1 2 1 1 ...
 $ q1length : num 14 10 14 11 13 17 4 7 8 9 ...
 $ q2length : num 12 15 10 13 7 17 11 9 7 9 ...
 $ diff_length : num 2 5 4 2 6 0 7 2 1 0 ...
 $ dist : num 9 41 30 39 41 15 51 15 5 24 ...
 $ jw_meth : num 0.0462 0.2791 0.2745 0.3855 0.3406 ...
 $ cosine_meth : num 0.0164 0.0644 0.0486 0.2052 0.0996 ...
 $ simi : num 0.862 0.506 0.444 0.203 0.301 ...
 $ lcs : num 9 47 60 69 61 45 59 27 5 50 ...
 $ dl : num 9 42 40 47 51 38 52 19 4 46 ...
 $ jaccard : num 0 0.348 0.261 0.409 0.333 ...
 $ osa : num 9 42 40 47 51 38 52 19 4 46 ...
```

```
> summary(text_process)
```

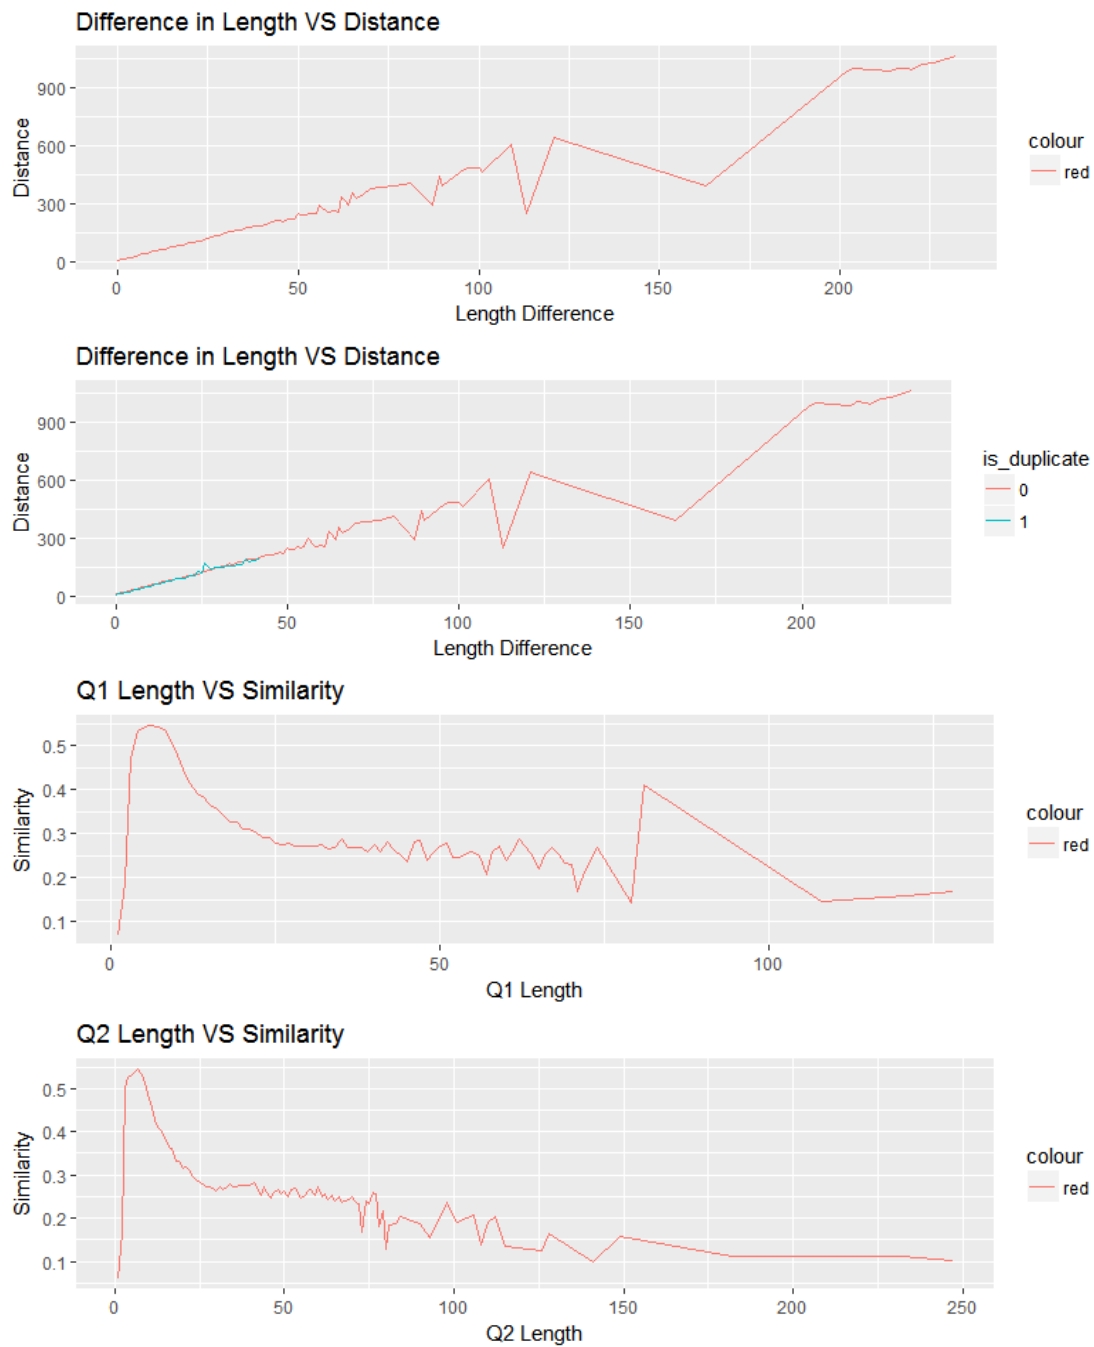
question1	question2	is_duplicate	q1length
Length:404274	Length:404274	0:255011	Min. : 1.00
Class :character	Class :character	1:149263	1st Qu.: 7.00
Mode :character	Mode :character		Median : 10.00
			Mean : 11.13
			3rd Qu.: 13.00
			Max. :128.00

q2length	diff_length	dist	jw_meth
Min. : 1.00	Min. : 0.00	Min. : 0.00	Min. :0.0000
1st Qu.: 7.00	1st Qu.: 1.00	1st Qu.: 14.00	1st Qu.:0.2063
Median : 10.00	Median : 2.00	Median : 23.00	Median :0.2779
Mean : 11.38	Mean : 3.78	Mean : 29.22	Mean :0.2657
3rd Qu.: 13.00	3rd Qu.: 5.00	3rd Qu.: 37.00	3rd Qu.:0.3342
Max. :247.00	Max. :232.00	Max. :1068.00	Max. :1.0000

cosine_meth	simi	lcs
Min. :0.00000	Min. :0.0000	Min. : 0.00
1st Qu.:0.03814	1st Qu.:0.3000	1st Qu.: 19.00
Median :0.06571	Median :0.4500	Median : 37.00
Mean :0.07792	Mean :0.4904	Mean : 47.42
3rd Qu.:0.10359	3rd Qu.:0.6552	3rd Qu.: 65.00
Max. :1.00000	Max. :1.0000	Max. :1076.00

dl	jaccard	osa
Min. : 0.00	Min. :0.0000	Min. : 0.00
1st Qu.: 15.00	1st Qu.:0.1053	1st Qu.: 15.00
Median : 30.00	Median :0.1875	Median : 30.00
Mean : 37.79	Mean :0.1943	Mean : 37.81
3rd Qu.: 51.00	3rd Qu.:0.2727	3rd Qu.: 51.00
Max. :1071.00	Max. :1.0000	Max. :1071.00

diff_length is created by subtracting the length of Question1 & Question2 with absolute values(non negative)



Here it shows that as the length of questions increases their similarity decreases. The same trend is also shown in difference in length(diff_length) below

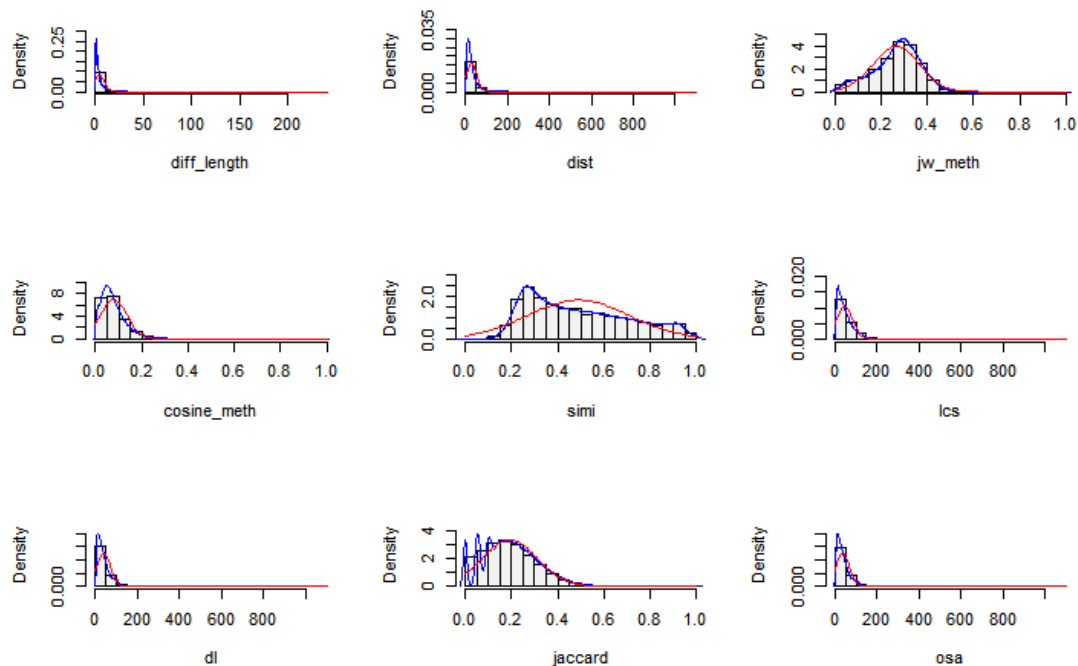


This shows as the difference in string length increases string similarity decreases.

```
> table(text_process$is_duplicate[which(text_process$q2length>20)])
  0      1
28172 4677
> table(text_process$is_duplicate[which(text_process$diff_length>1)])
  0      1
168682 78389
> sum(text_process$is_duplicate==0 & text_process$diff_length>10)/nrow(text_process)*100
[1] 7.457813
> sum(text_process$is_duplicate==1 & text_process$diff_length>10)/nrow(text_process)*100
[1] 0.7724959
```

This gives us a clue that as the difference in string length increases more than 10 there is a high chance that they are not same. We can approve it after applying a model.

A combined plot for all numeric data



Here dist, dl and osa showing the same distribution. So we will check for multicollinearity and remove unnecessary variables.

Checking for Multicollinearity

```
> vif(as.data.frame(t_matrix))
  Variables      VIF
1  q1length  2.001873
2  q2length  2.952932
3 diff_length  6.362810
4      dist  31.394486
5   jw_meth  5.945232
6 cosine_meth  4.057238
7      simi  10.331273
8       lcs  57.094787
9       dl 35008.220338
10    jaccard  3.667862
11      osa 35136.567778
> vifcor(t_matrix[,1:11], th=0.95)
2 variables from the 11 input variables have collinearity problem:
osa dl
```

After excluding the collinear variables, the linear correlation coefficients ranges between:
min correlation (jaccard ~ q2length): -0.001888916
max correlation (lcs ~ dist): 0.9118671

```
----- VIFs of the remained variables -----
  Variables      VIF
1  q1length  2.085014
2  q2length  2.689762
3 diff_length  5.066233
4      dist 18.135242
5   jw_meth  6.183475
6 cosine_meth  3.358180
7      simi  8.496280
8       lcs 18.322263
9    jaccard  3.448290
```

From this we see that dl & osa show the maximum collinearity so they are not fed to the model.

Word Cloud

Question 1 Word Cloud



Question 2 Word Cloud



A decision tree **C5.0** is used for classification with accuracy 71%

```
> confusionMatrix(xtab)
```

Confusion Matrix and Statistics

	predicted	
observed	0	1
0	100140	28881
1	28750	46511

Accuracy : 0.7179
95% CI : (0.7159, 0.7198)
No Information Rate : 0.6309
P-Value [Acc > NIR] : <2e-16

Kappa : 0.394
McNemar's Test P-Value : 0.5881

Sensitivity : 0.7769
Specificity : 0.6169
Pos Pred Value : 0.7762
Neg Pred Value : 0.6180
Prevalence : 0.6309
Detection Rate : 0.4902
Detection Prevalence : 0.6316
Balanced Accuracy : 0.6969

'Positive' Class : 0

Another decision tree “**rpart**” is used for classification with accuracy 68%

```
> confusionMatrix(xtab)
```

Confusion Matrix and Statistics

	predicted	
observed	0	1
0	107073	21948
1	42105	33156

Accuracy : 0.6864
95% CI : (0.6844, 0.6885)
No Information Rate : 0.7303
P-Value [Acc > NIR] : 1

Kappa : 0.2864
McNemar's Test P-Value : <2e-16

Sensitivity : 0.7178
Specificity : 0.6017
Pos Pred Value : 0.8299
Neg Pred Value : 0.4405
Prevalence : 0.7303
Detection Rate : 0.5241
Detection Prevalence : 0.6316
Balanced Accuracy : 0.6597

'Positive' Class : 0

Random Forest Model

A Random Forest model is used to predict "is_duplicate" on test data with 85% accuracy.

```
> confusionMatrix(xtab)
```

Confusion Matrix and Statistics

```
      predicted
observed 0      1
0 115095 13926
1  16401 58860

      Accuracy : 0.8515
      95% CI : (0.85, 0.8531)
No Information Rate : 0.6437
P-Value [Acc > NIR] : < 2.2e-16

      Kappa : 0.6788
McNemar's Test P-Value : < 2.2e-16

      Sensitivity : 0.8753
      Specificity : 0.8087
      Pos Pred Value : 0.8921
      Neg Pred Value : 0.7821
      Prevalence : 0.6437
      Detection Rate : 0.5634
      Detection Prevalence : 0.6316
      Balanced Accuracy : 0.8420

      'Positive' Class : 0
```

```
> rf_model
```

Call:

```
randomForest(formula = factor(is_duplicate) ~ diff_length + dist +
jw_meth + cosine_meth + simi + lcs + jaccard, data = train,
importance = TRUE, ntree = 100)
```

Type of random forest: classification

Number of trees: 100

No. of variables tried at each split: 2

OOB estimate of error rate: 29.31%

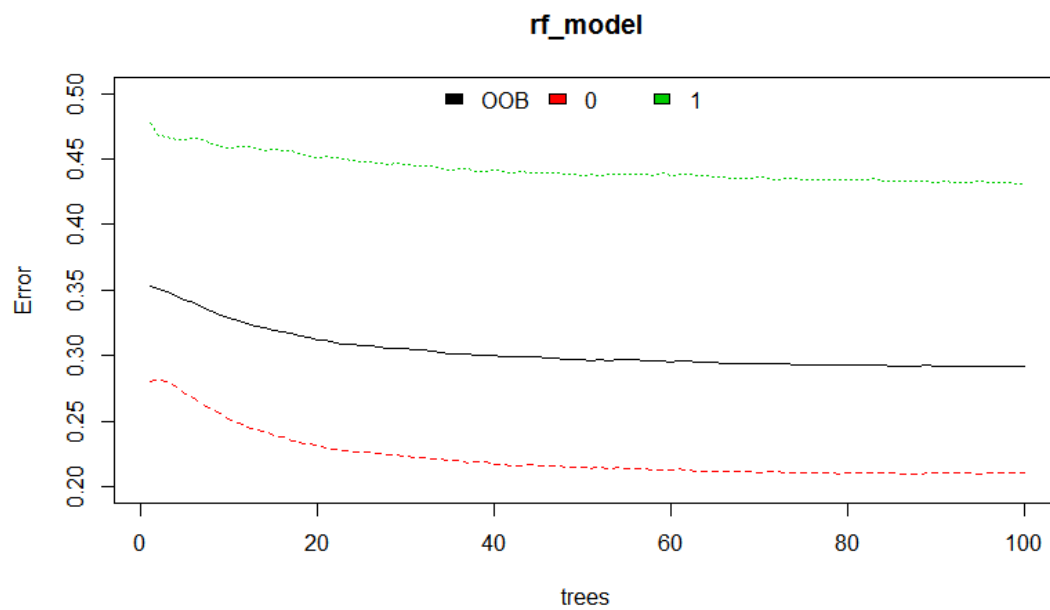
Confusion matrix:

```
      0      1 class.error
0 98981 26979  0.2141870
1 31642 42398  0.4273636
```


It is noted that as sample is increased or if tree is increased there is decrease in error of class 1.

It means as the model trains on more data it is better able to predict for class 1

Error Plot



Variable Importance



The above variable importance plot shows that the methods stringsim(sim), Jaro distance(jw_meth), Cosine distance(cosine_meth) are important predictors to the model.

Now we will look at the predicted and actual values from RF Model

test_model variable is created which contains the observations in which observed class is same as predicted class.

```
> table(test_model$is_duplicate)

      0      1
114598 59393

> sum(test_model$is_duplicate==0 & test_model$diff_length>10)/nrow(test_model)*100
[1] 8.720566
> sum(test_model$is_duplicate==1 & test_model$diff_length>10)/nrow(test_model)*100
[1] 0.5483042
```

The above code shows that what we have earlier seen in the train data is correctly predicted in test data that as the **difference in string length increases more than 10 there is a high chance that they are not same.**

Model Evaluation

Among the decision tree models Random Forest gave us the maximum accuracy. But still predicting for string similarity is not easy.

KNN Model

KNN is used with K=1 with 96% accuracy

```
> Conf_matrix
```

```
pred      0      1
      0 125917  4392
      1   2669 71305
```

```
>
```

```
> accuracy = sum(diag(Conf_matrix))/nrow(test)
```

```
> accuracy
```

```
[1] 0.9654352
```

K=3 with 95% accuracy

```
> Conf_matrix
```

```
pred      0      1
      0 126019  7670
      1   2485 68110
```

```
>
```

```
> accuracy = sum(diag(Conf_matrix))/nrow(test)
```

```
> accuracy
```

```
[1] 0.9502898
```

K=5 with 95% accuracy

```
> Conf_matrix
```

```
pred      0      1
      0 126280  8958
      1   2224 66822
```

```
>
```

```
> accuracy = sum(diag(Conf_matrix))/nrow(test)
```

```
> accuracy
```

```
[1] 0.9452625
```

Model Evaluation

K-Nearest neighbour classification is a general technique to learn classification based on instance and do not have to develop an abstract model from the training data set. However the classification process could be very expensive because it needs to compute the similarity values individually between the test and training examples. Choosing a proper K value is important to the performance of K-Nearest Neighbour classifier. If k value is too large the nearest neighbour classifier may misclassify the test instance because its list of nearest neighbours may include data points that are located far away from its neighbourhood. On the other hand, if k value is too small, then the nearest neighbour classifier may be susceptible to over fitting because of noise in the training data set.